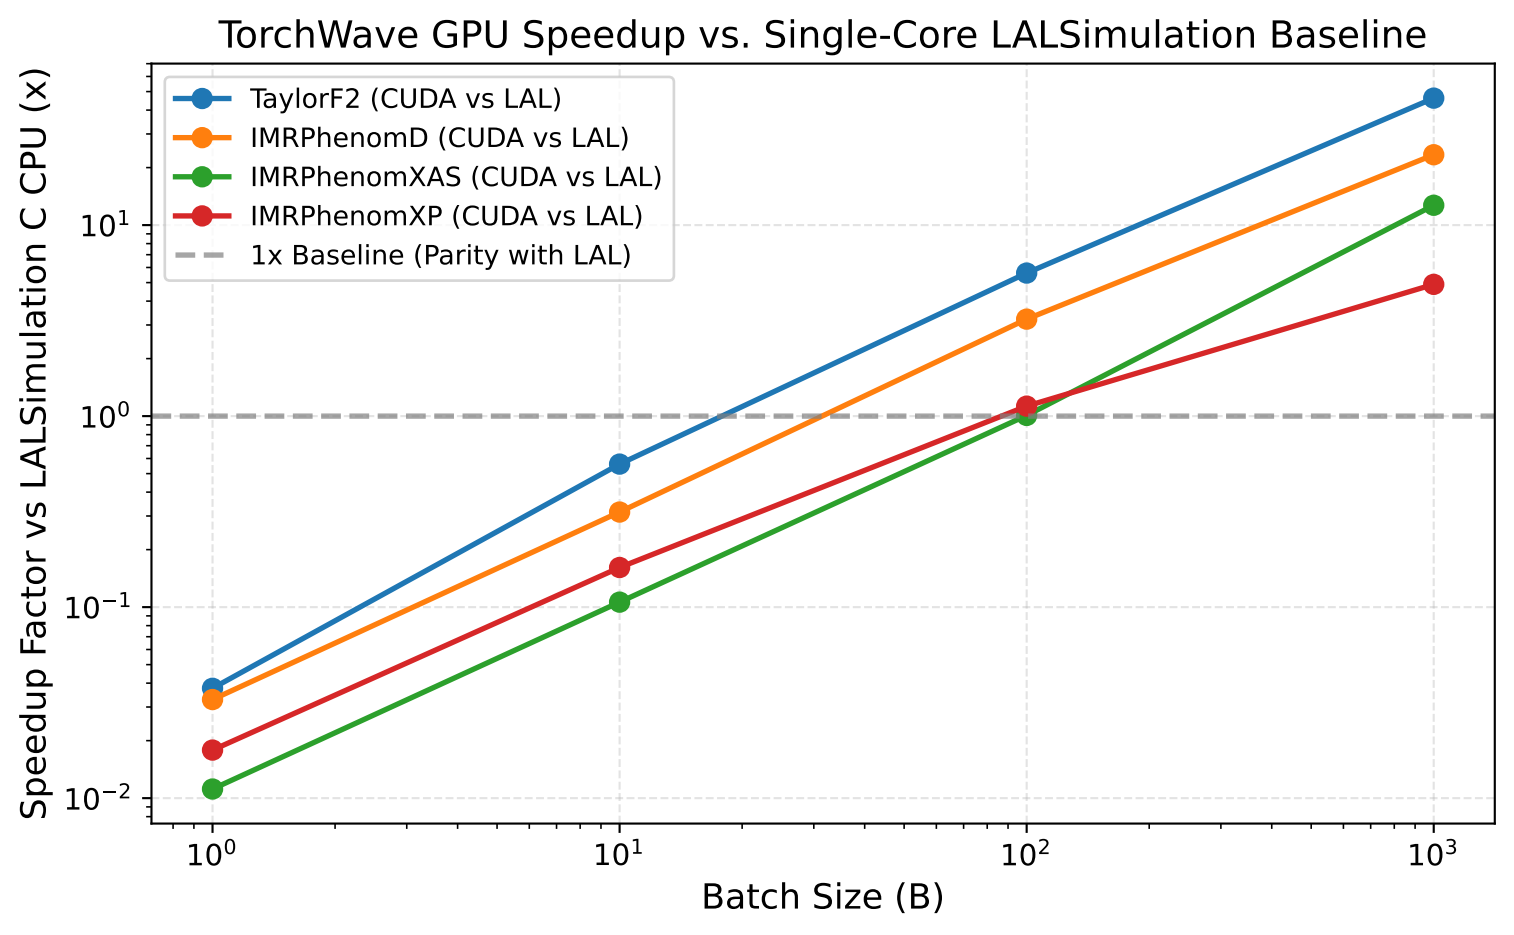
TorchWave GPU Speedup vs. Single-Core LALSimulation Baseline
TaylorF2 (CUDA vs LAL)
IMRPhenomD (CUDA vs LAL)
IMRPhenomXAS (CUDA vs LAL)
IMRPhenomXP (CUDA vs LAL)
1x Baseline (Parity with LAL)
Speedup Factor vs LALSimulation C CPU (x)
Batch Size (B)
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{0}
10^{1}
10^{2}
10^{3}